


GAN


- Generative adversarial network
- Generative model
- Map a latent code to high dimension data
- Training in an adversarial way

GAN Architecture

Noise vector




Generated data




GAN Architecture

Noise vector



Generated data




Real data




GAN Architecture


Noise vector



Generated data



Real data




Discriminator



GAN Architecture

Noise vector



Generated data



Real data




Discriminator




GAN Architecture


Noise vector



Generated data



Real data




Discriminator



GAN Architecture

Noise vector



Generated data



Real data




Discriminator



GAN Architecture

Noise vector



Generated data



Real data



Discriminator



GAN Architecture

Noise vector


z



Generated data



Real data



Discriminator

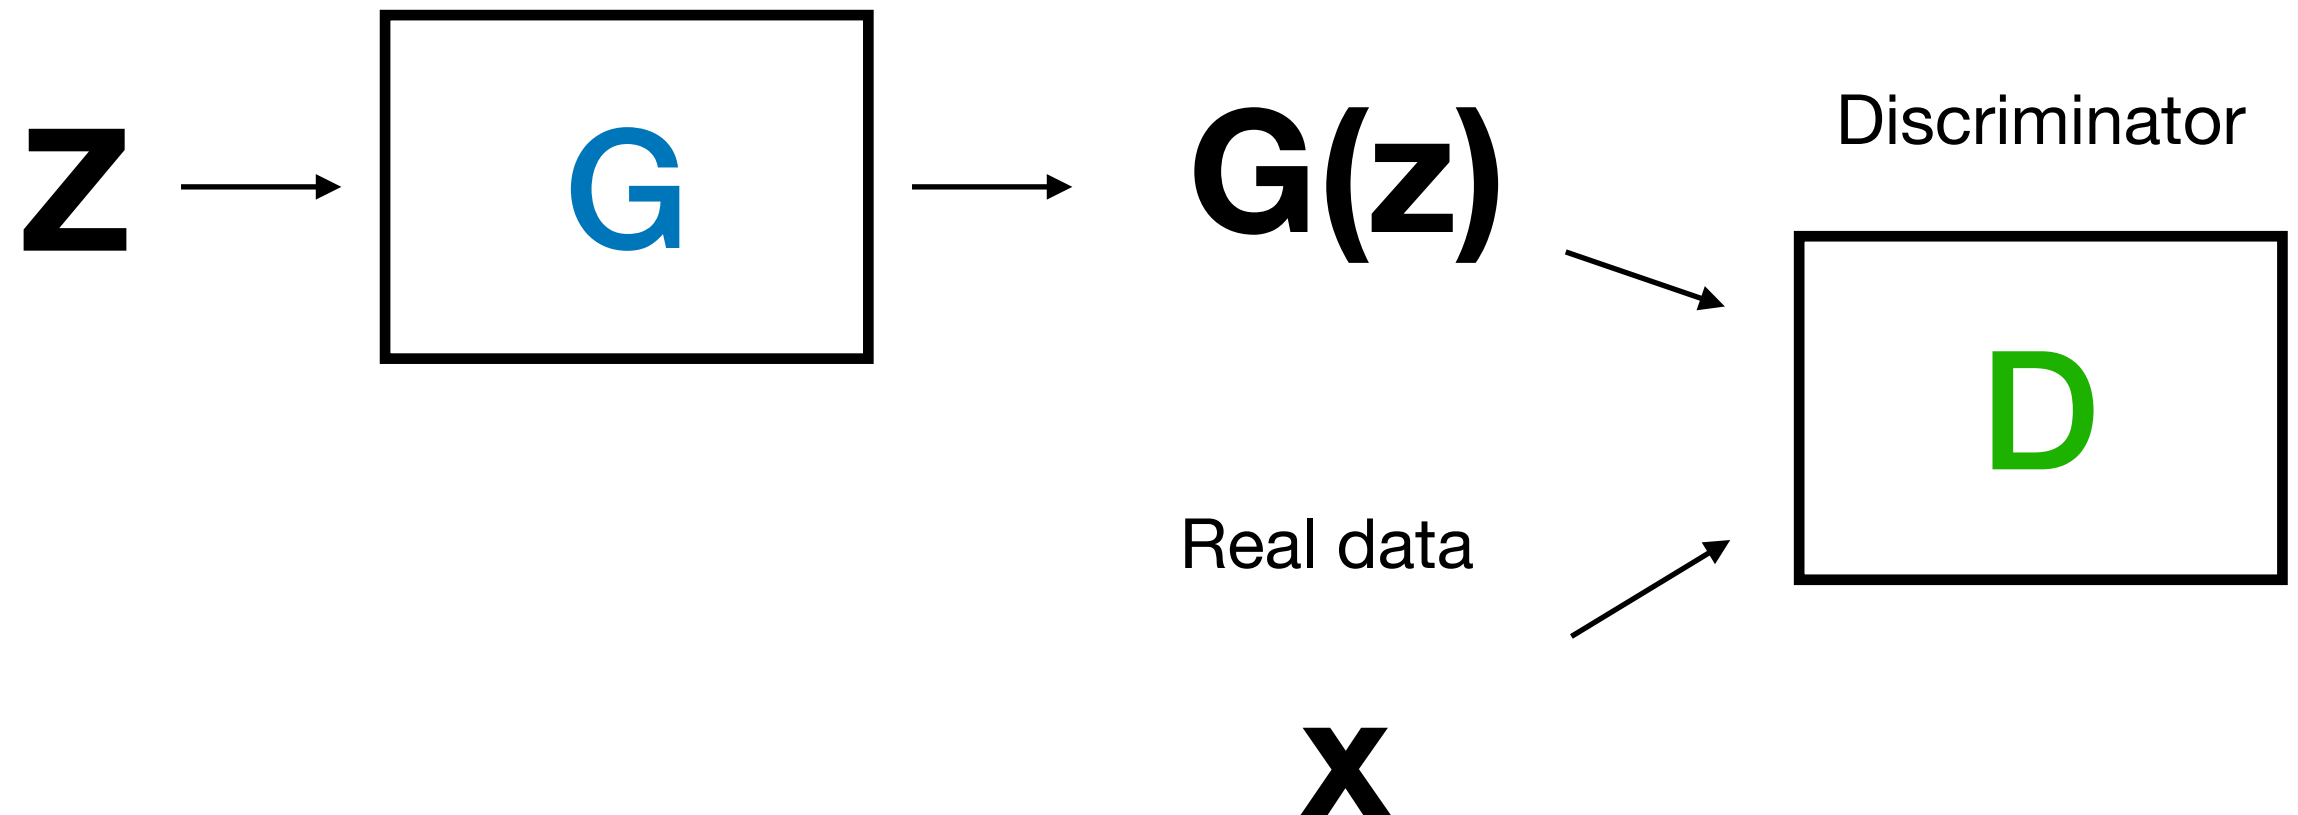


- Noise vector z
- Generator **G** maps from noise z to data space $G(z)$
- Discriminator **D** distinguishes $G(z)$ and real data x

GAN Architecture


Noise vector

Generated data




- Noise vector \mathbf{z}
- Generator G maps from noise \mathbf{z} to data space $G(\mathbf{z})$
- Discriminator D distinguishes $G(\mathbf{z})$ and real data \mathbf{x}

Minimax Game



- Train discriminator \mathbf{D} to **maximise** the probability of assigning the correct label to both training samples \mathbf{x} (label as 1, i.e. $\mathbf{D}(\mathbf{x})$: $\mathbf{1}$) and samples from \mathbf{G} (label as 0, i.e. $\mathbf{D}(\mathbf{G}(\mathbf{z}))$: $\mathbf{0}$)
- Train generator \mathbf{G} to **minimise** the distance between the distribution of generated samples $\mathbf{G}(\mathbf{z})$ and the distribution of training samples \mathbf{x}


Minimax Game



Full objective:

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$


Global optimum



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

- For G , $p_g = p_{\text{data}}$
- For D , it outputs $D^*(x) = D^*(G(z))$

Loss function




$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

- For G , $\frac{1}{m} \sum_{i=1}^m \log \left(1 - D \left(G \left(\mathbf{z}^{(i)} \right) \right) \right)$

- For D , $\frac{1}{m} \sum_{i=1}^m \left[\log D \left(\mathbf{x}^{(i)} \right) + \log \left(1 - D \left(G \left(\mathbf{z}^{(i)} \right) \right) \right) \right]$


Loss function



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$


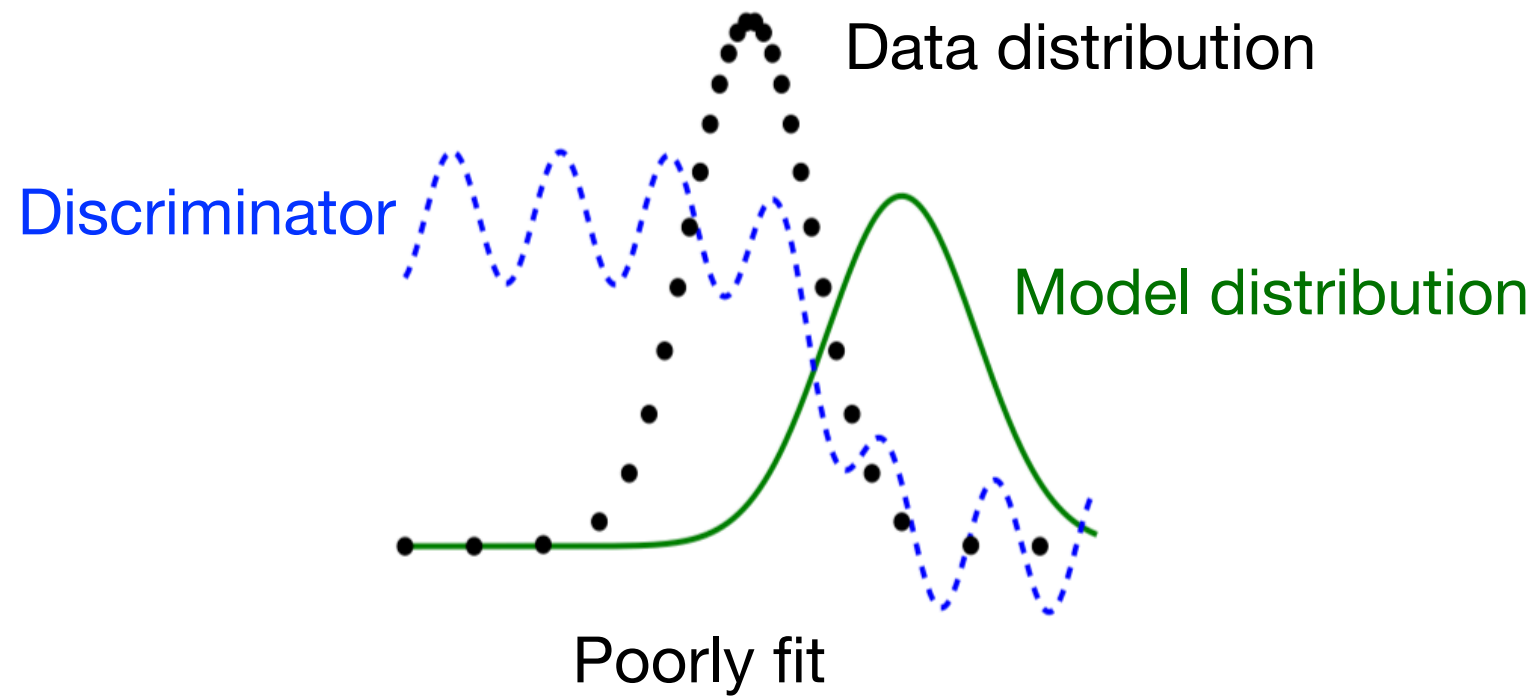
- For G , $\frac{1}{m} \sum_{i=1}^m \log \left(1 - D \left(G \left(\mathbf{z}^{(i)} \right) \right) \right)$
`tf.nn.sigmoid_cross_entropy_with_logits`
- For D , $\frac{1}{m} \sum_{i=1}^m \left[\log D \left(\mathbf{x}^{(i)} \right) + \log \left(1 - D \left(G \left(\mathbf{z}^{(i)} \right) \right) \right) \right]$

Training Procedure




- Use **SGD** (Stochastic gradient descent) or **Adam** optimiser on generator and discriminator
- Alternately update **G** and **D** until converge


GAN



After updating D



After updating G



Mixed strategy equilibrium

GAN Architecture

Noise vector



Generator




Generated data



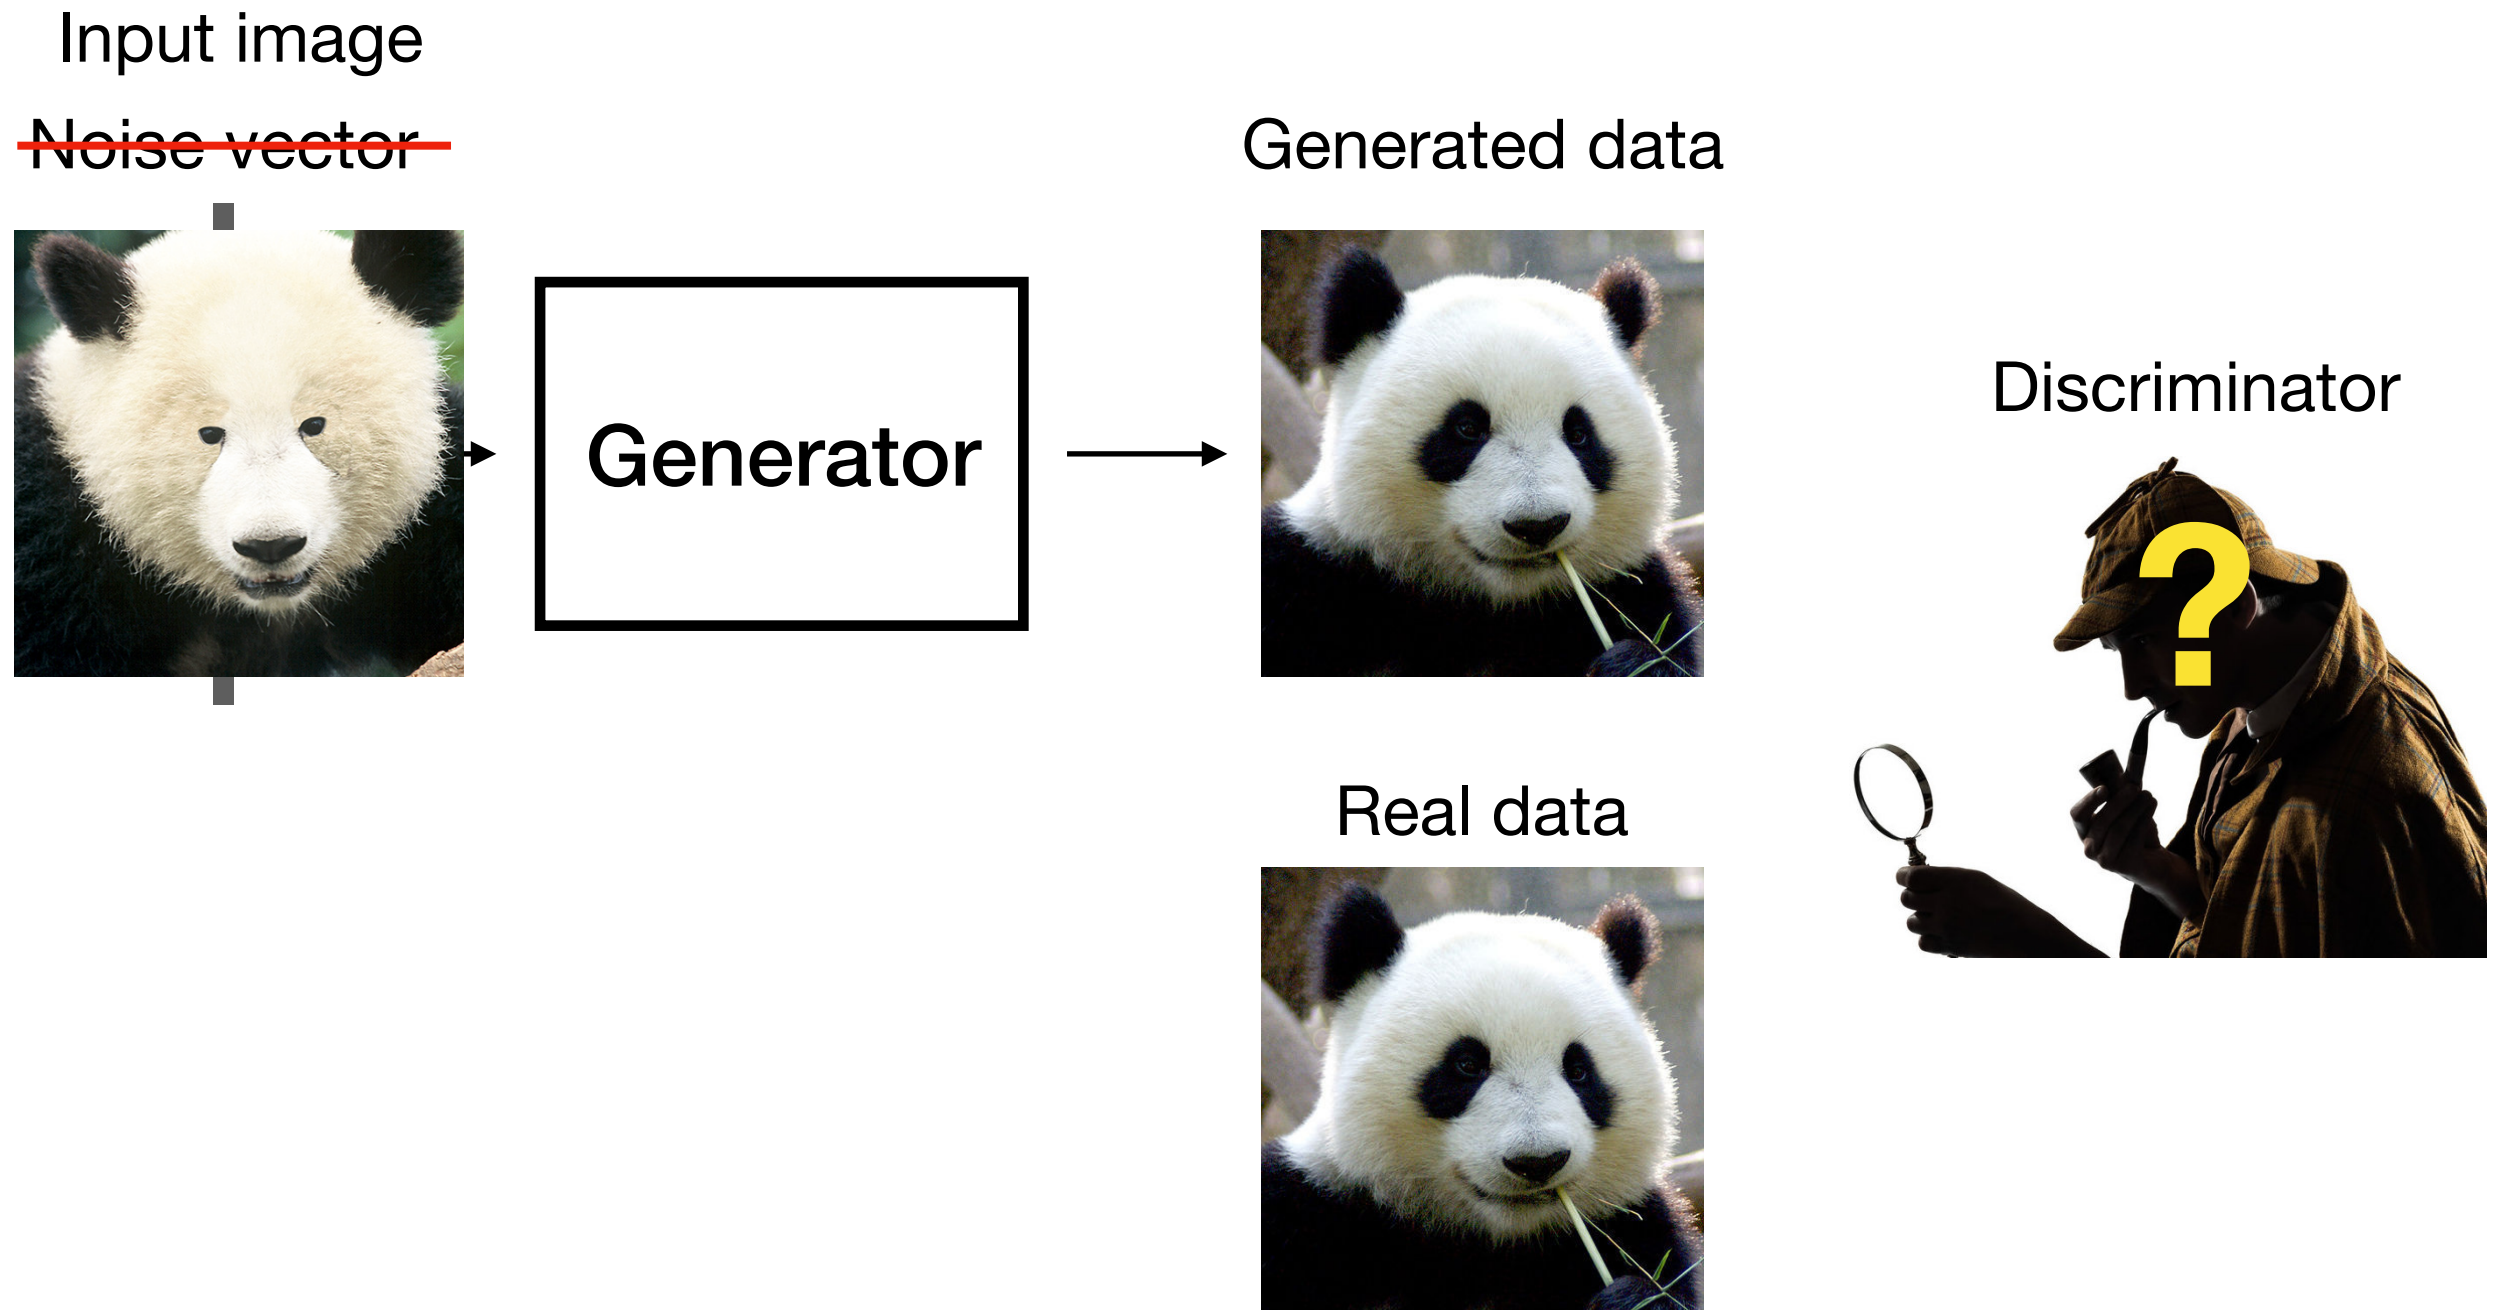
Real data



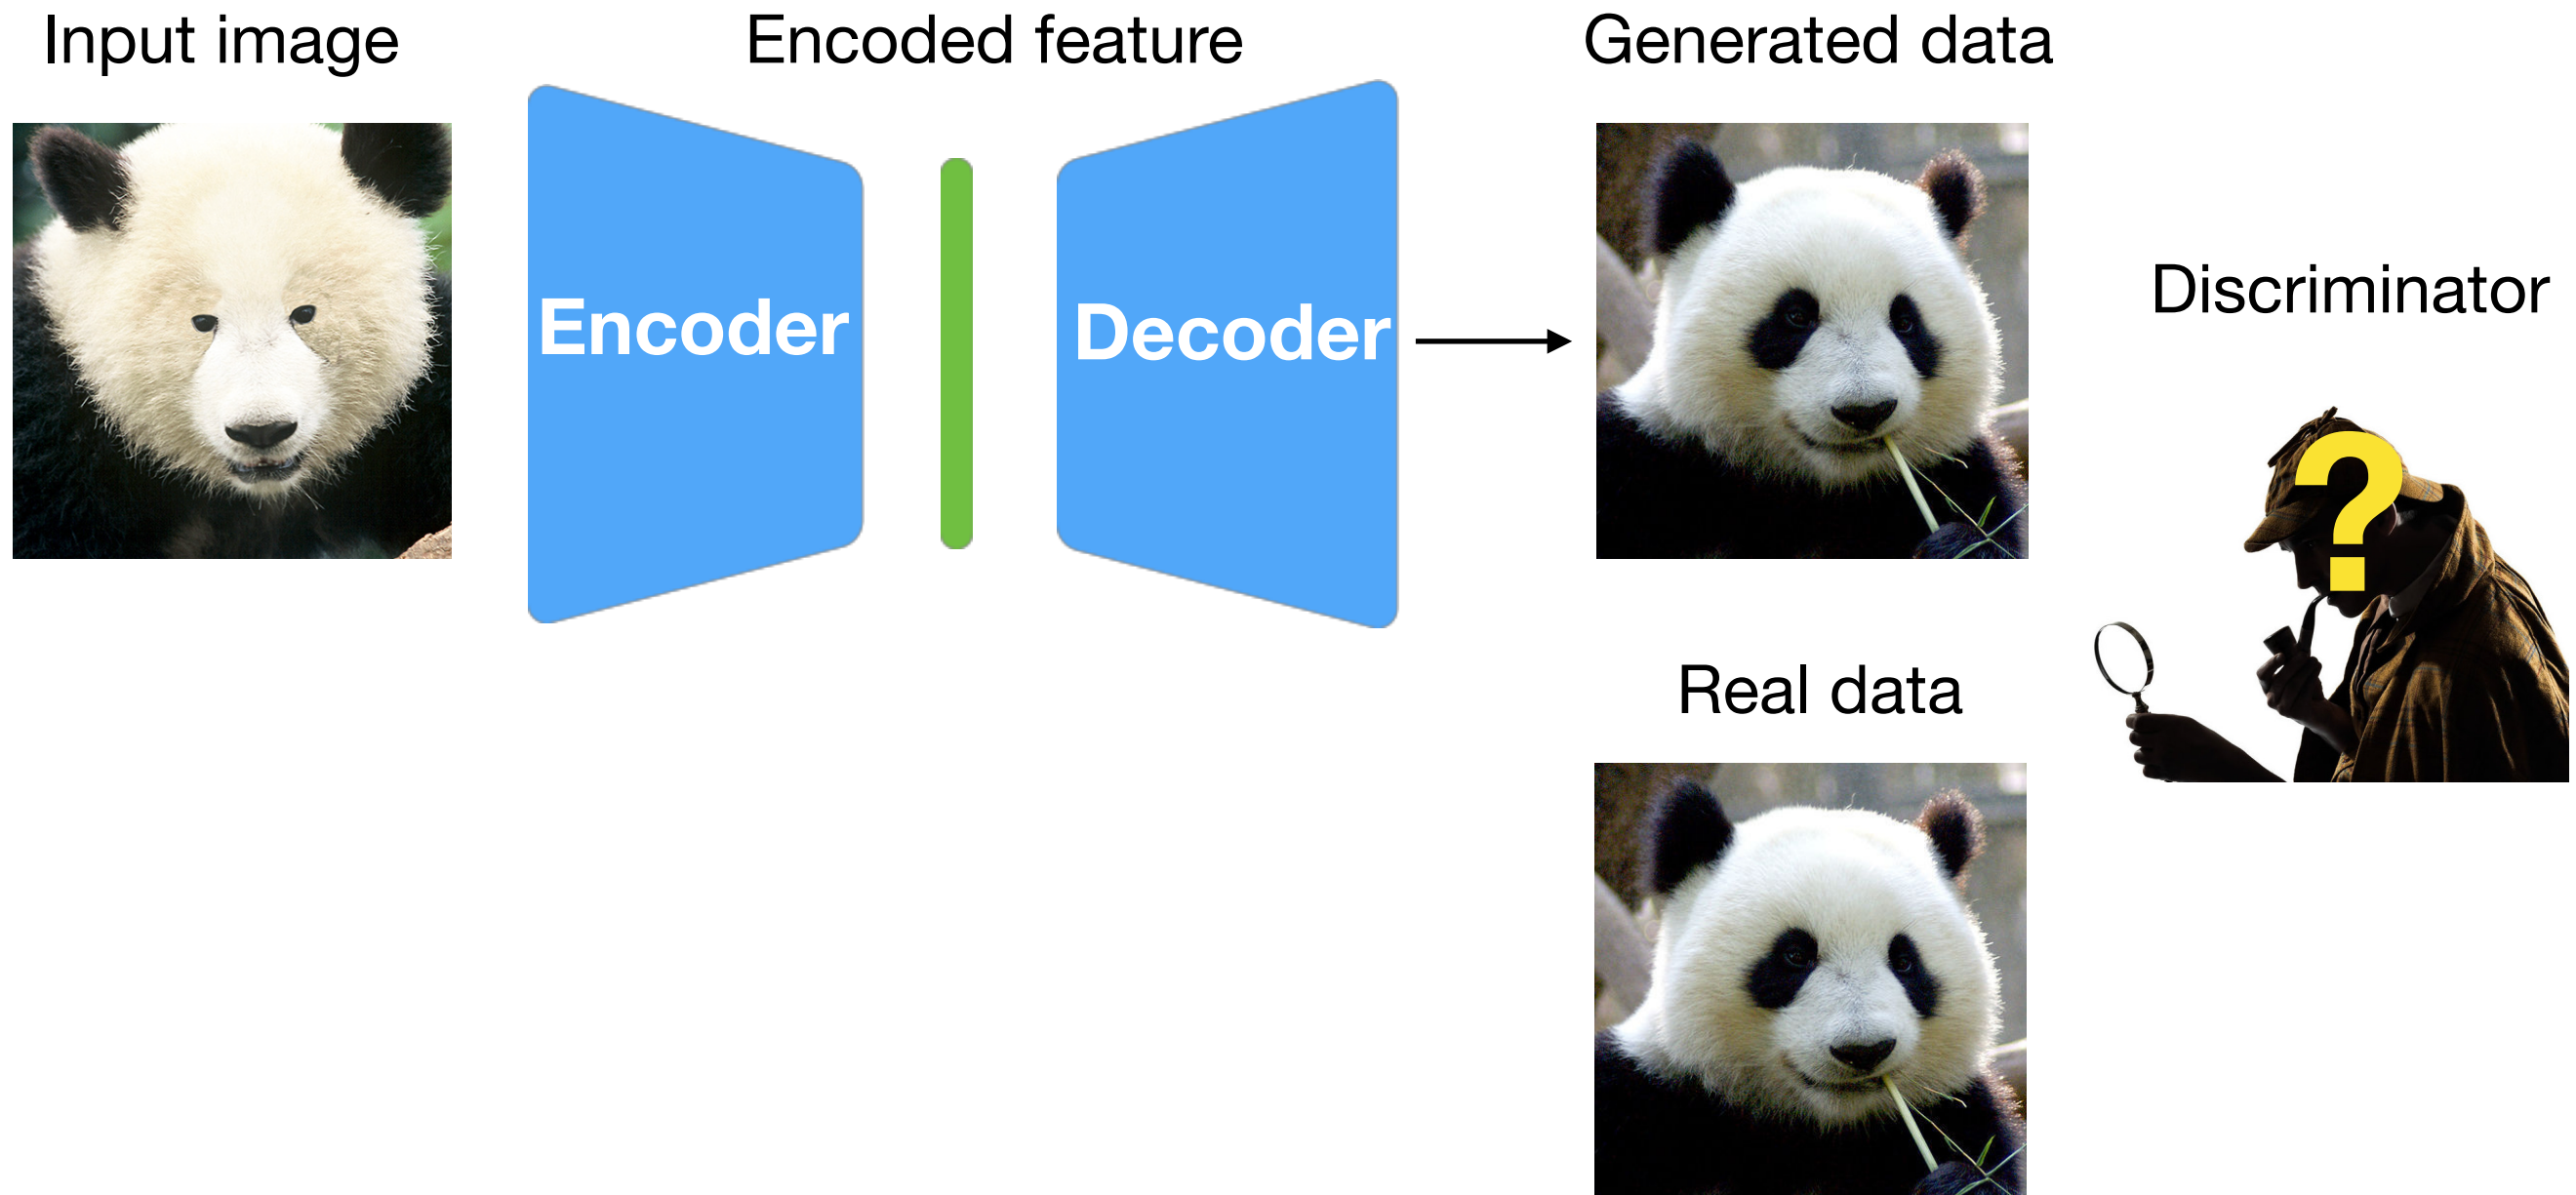
Discriminator



GAN Architecture



GAN with Autoencoder




Autoencoder

Input image



Encoded feature



Generated data




Autoencoder


Input image




Encoded feature



Generated data




Real data




L1
L2
...
Loss

Autoencoder


Input image



Encoded feature




Generated data



- Convolutional layer (`tf.nn.conv2d`). Set `stride>1` to downsample
- Convolutional layer + pooling
- fully connected layer (`tf.layers.dense`)

Autoencoder




- deconvolutional layer (`tf.nn.conv2d_transpose`). Set `stride>1` to upsample
- fully connected layer (`tf.layers.dense`) and reshape the output of last layer to the scale of the image

Disentangling

- Approach: Disentangling content and the rest (style, noise etc.) (supervise with weak label)


Input: pair of text image, same content different style



Disentangling

- Approach: Disentangling content and the rest (style, noise etc.)

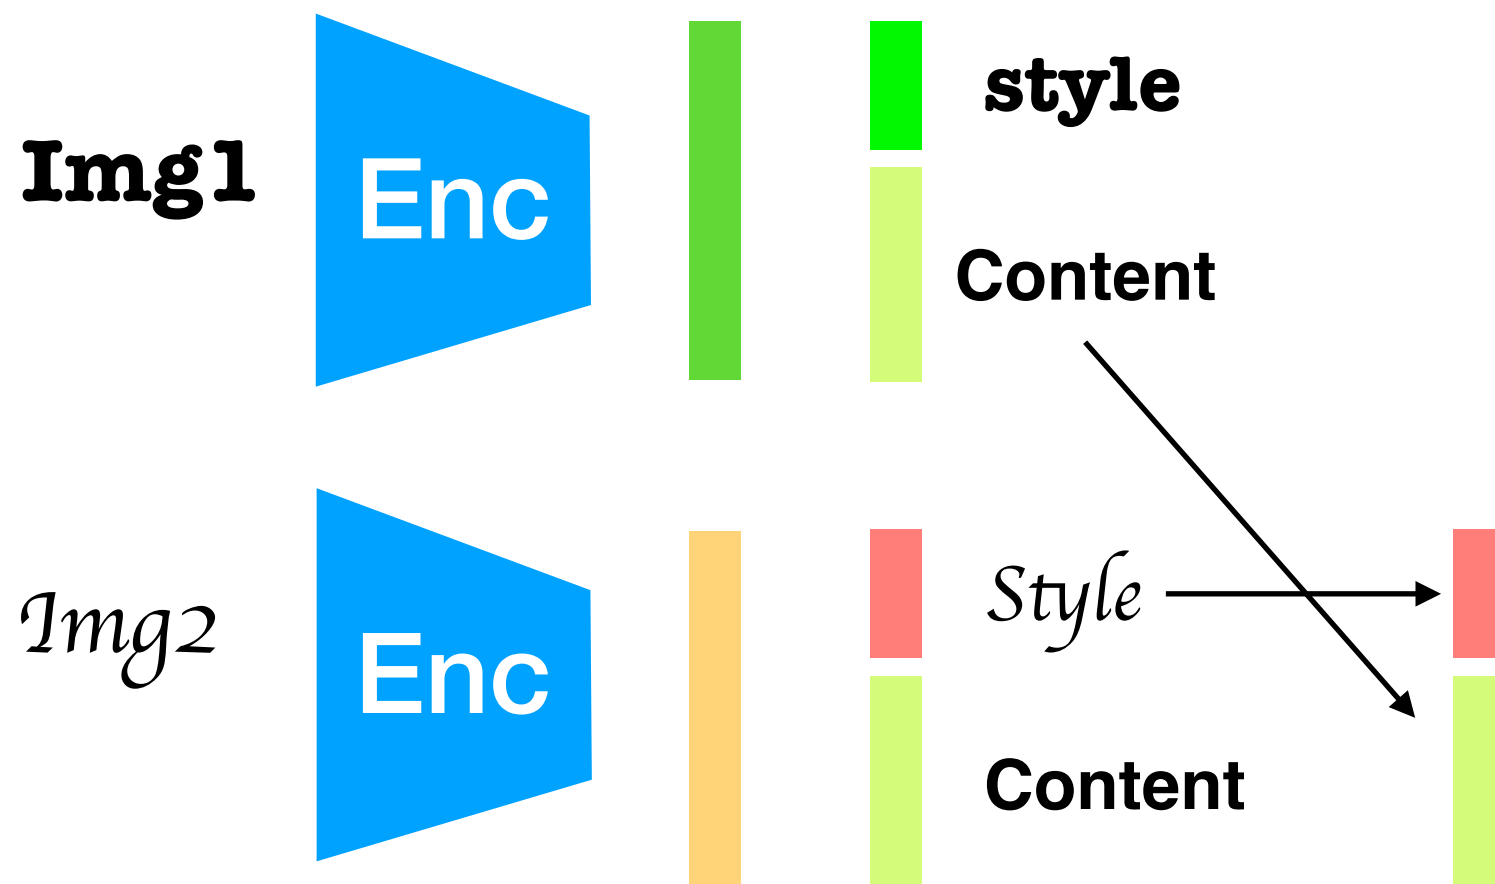
Input: pair of text image, same content different style



Disentangling

- Approach: Disentangling content and the rest (style, noise etc.)

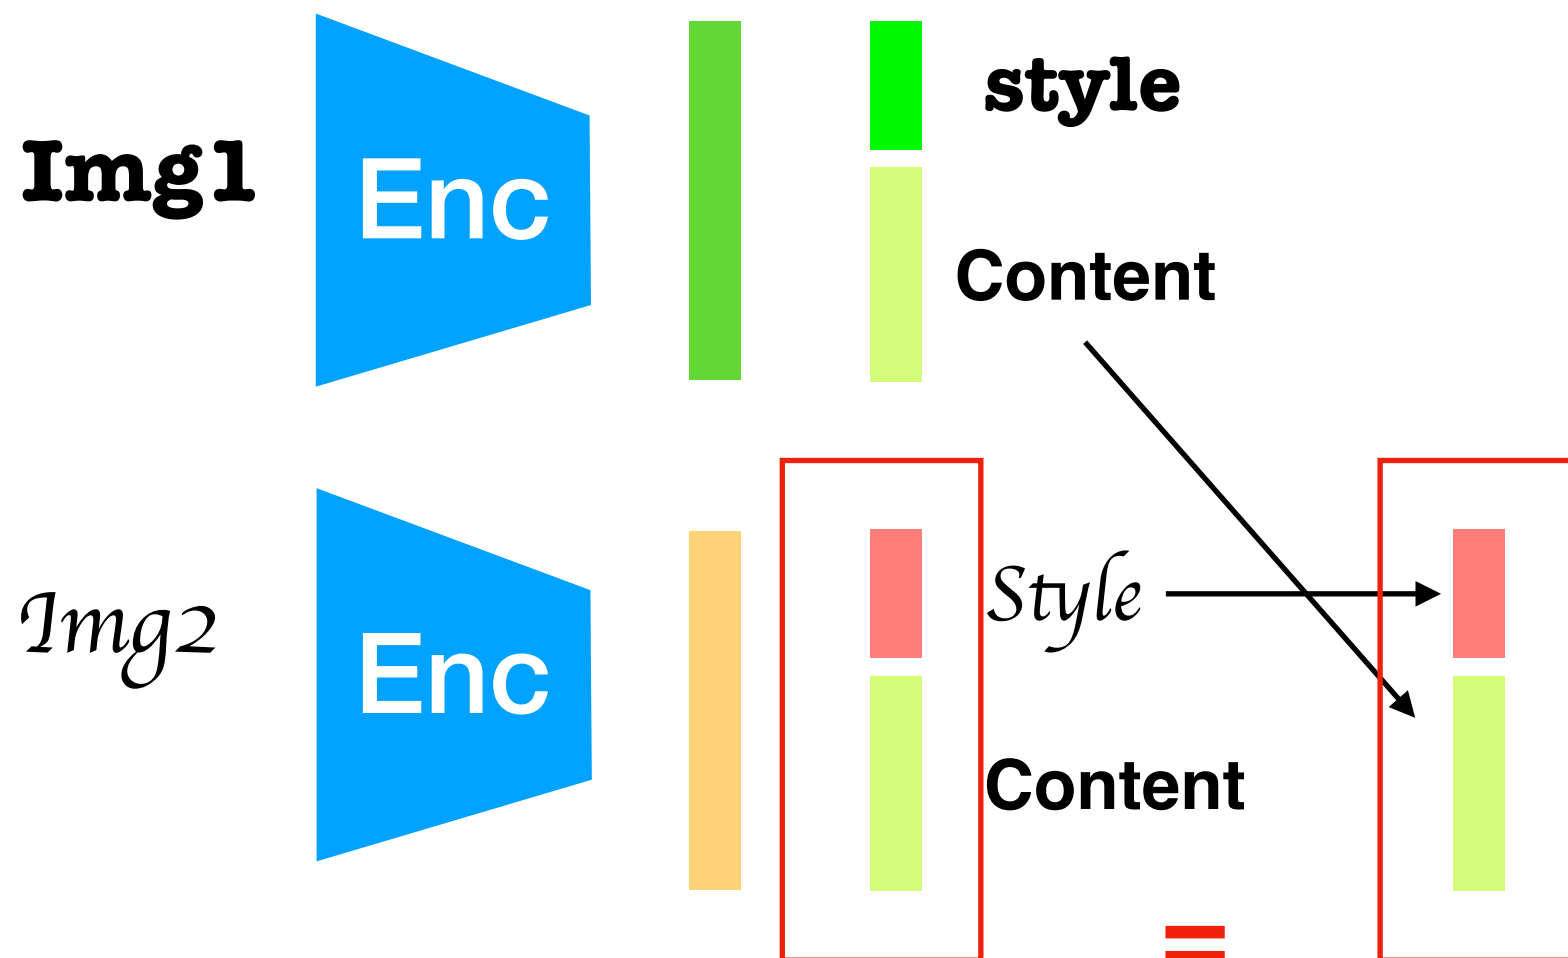
Input: pair of text image, same content different style



Disentangling

- Approach: Disentangling content and the rest (style, noise etc.)

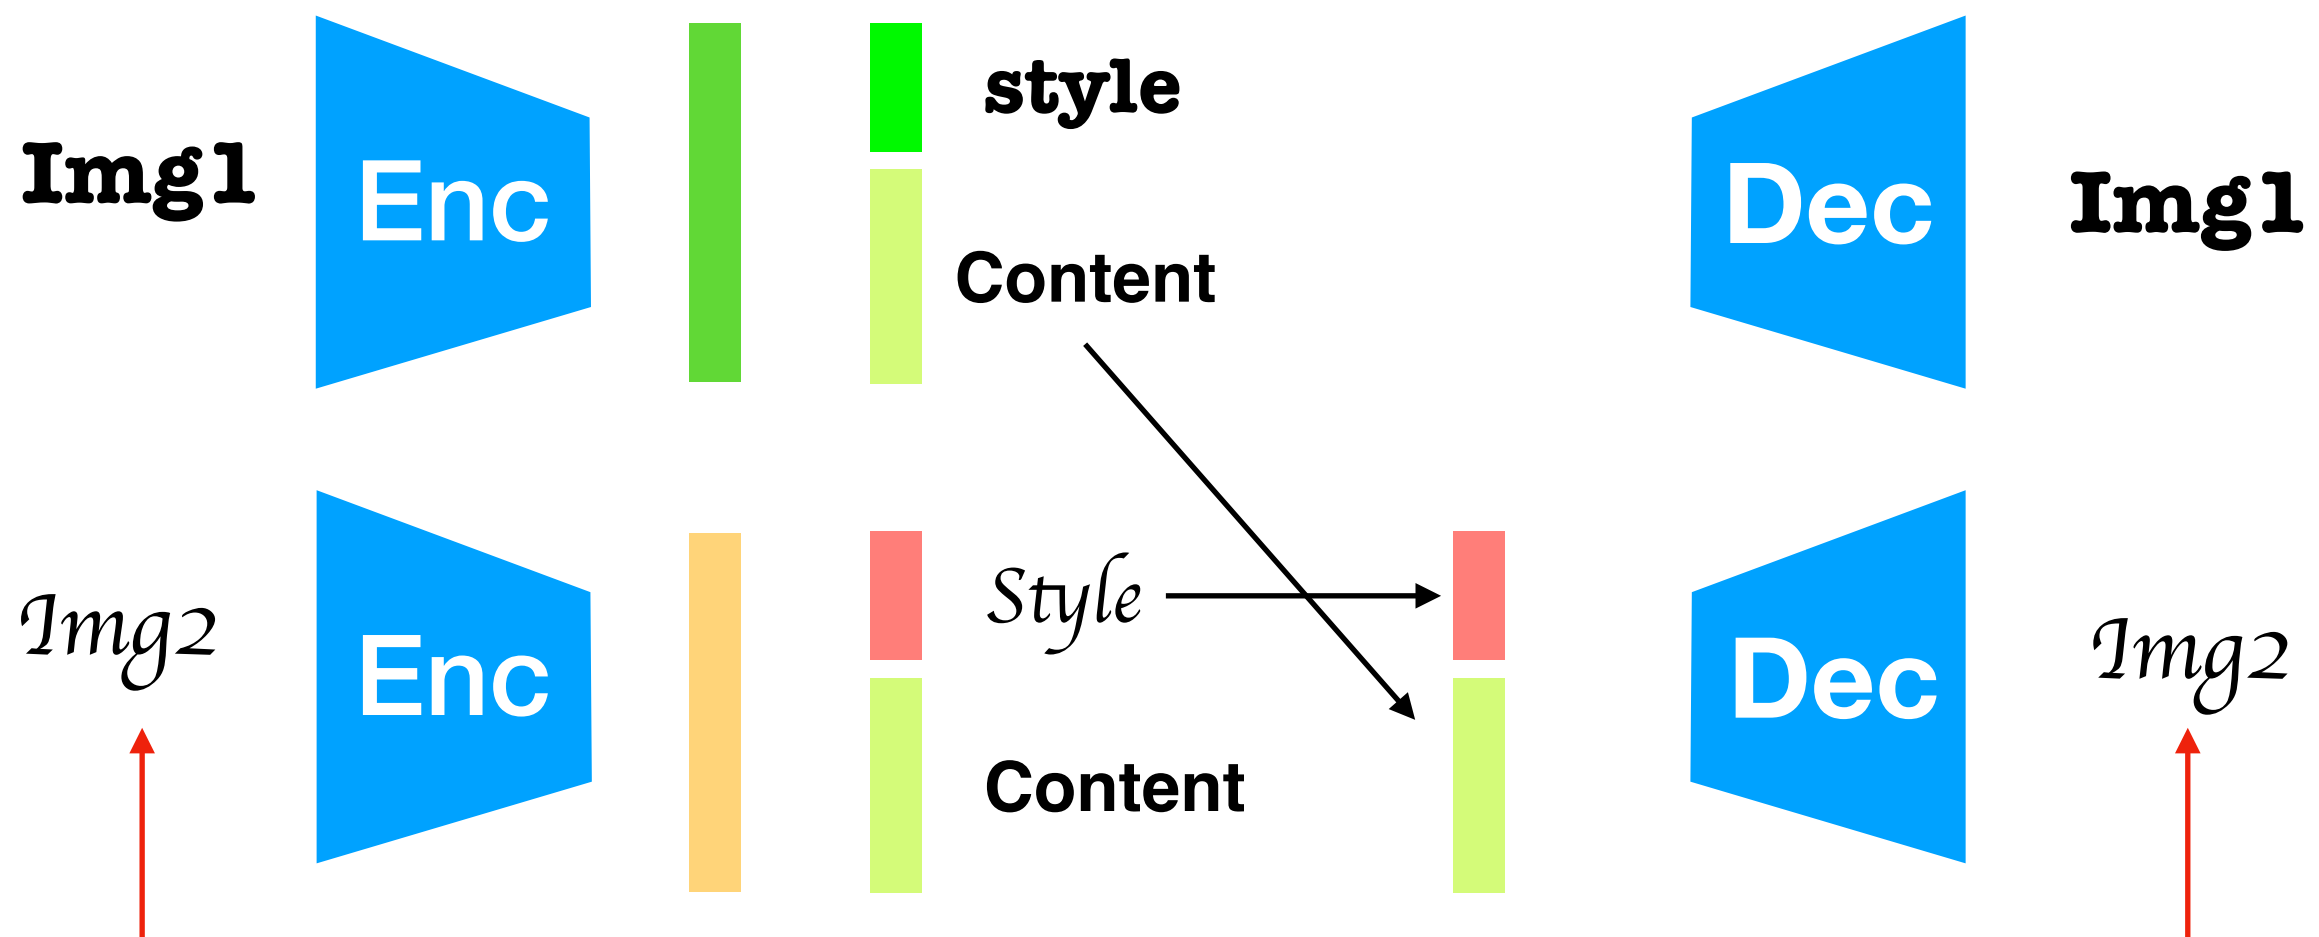
Input: pair of text image, same content different style



Disentangling

- Approach: Disentangling content and the rest (style, noise etc.)

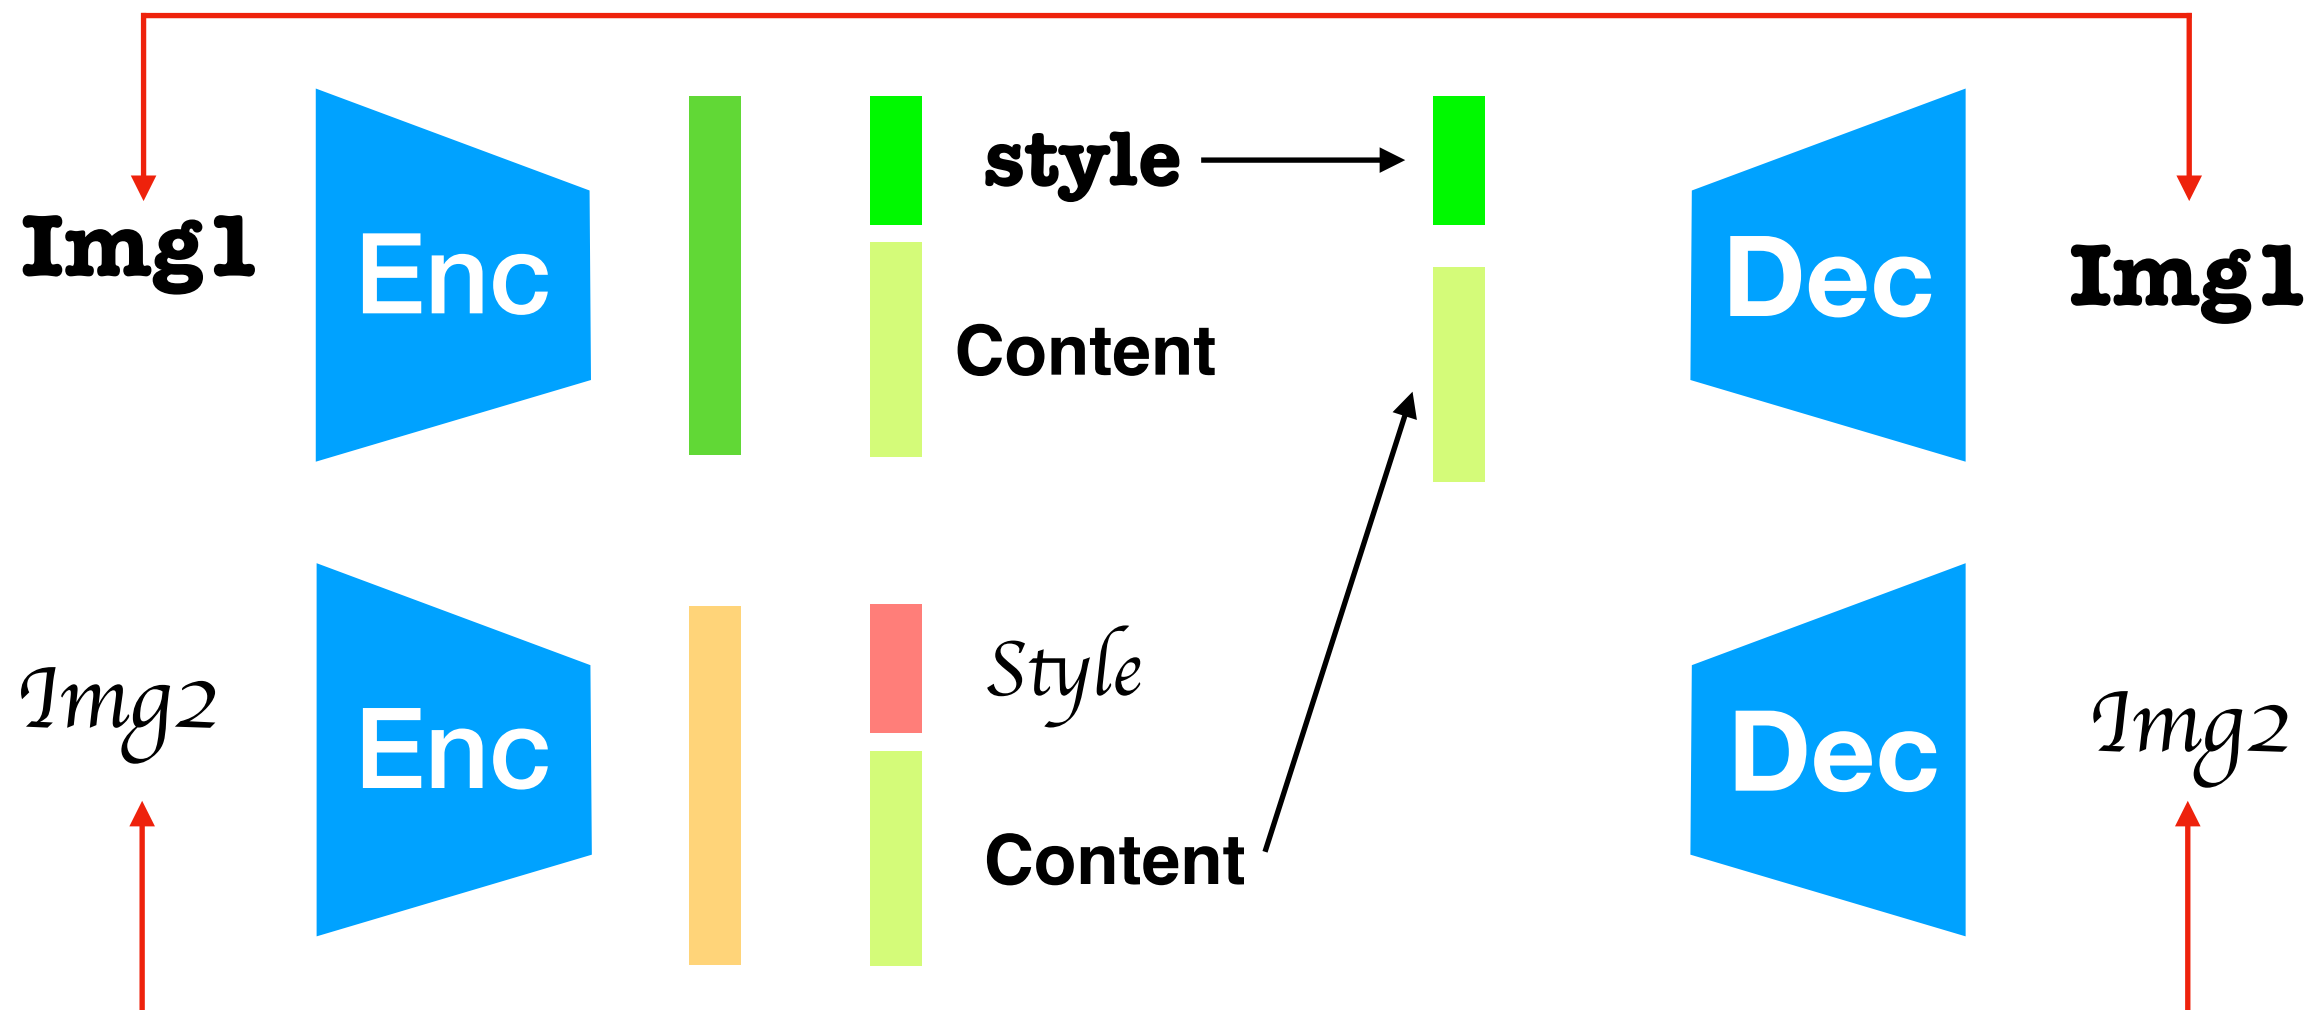
Input: pair of text image, same content different style



Disentangling

- Approach: Disentangling content and the rest (style, noise etc.)

Input: pair of text image, same content different style



Disentangling

- Approach: Disentangling content and the rest (style, noise etc.)

Output: Real data content + synthetic style

